An example of vectorization – Vector instructions improve the performance by processing multiple data items concurrently.

```
for (int i=0; i<16; i++)
    a[i] = b[i] + c[i];
```

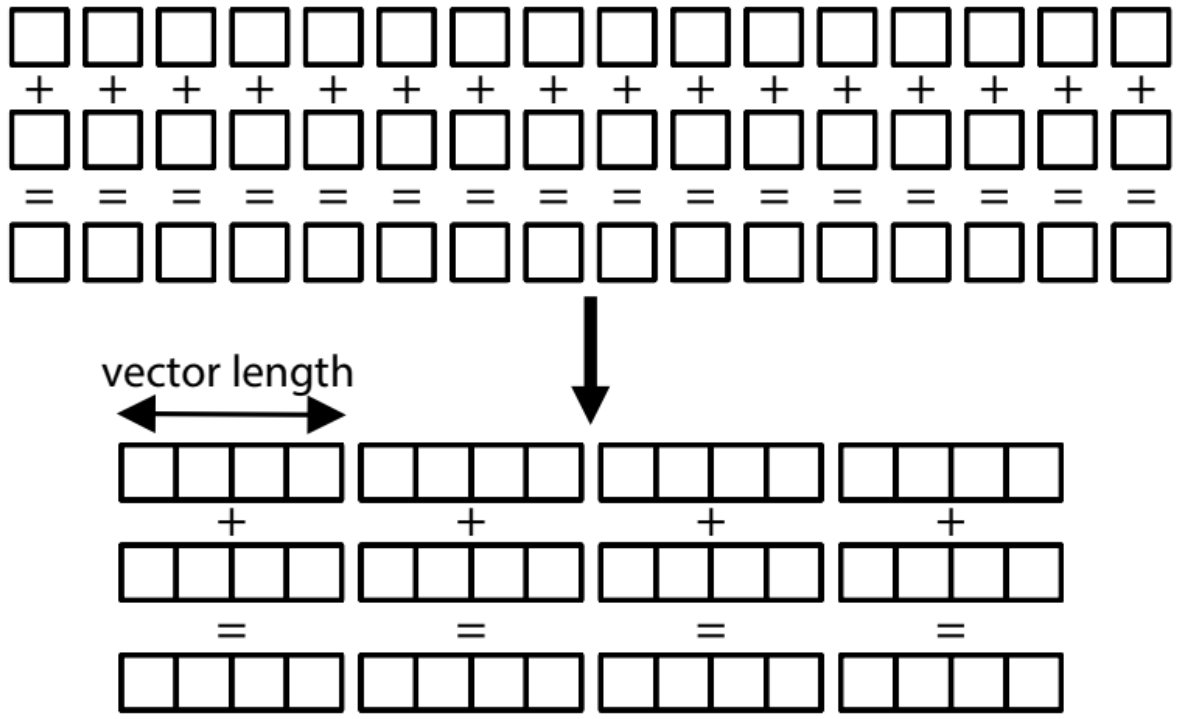

Scalar instructions

32 loads
16 adds
16 stores

SIMD instructions

8 loads
4 adds
4 stores

Loop modifications needed to vectorize a loop using SIMD instructions – To enforce alignment and ensure the remaining loop length is a multiple of the vector length, the compiler may need to peel off iterations and also treat the tail end separately.

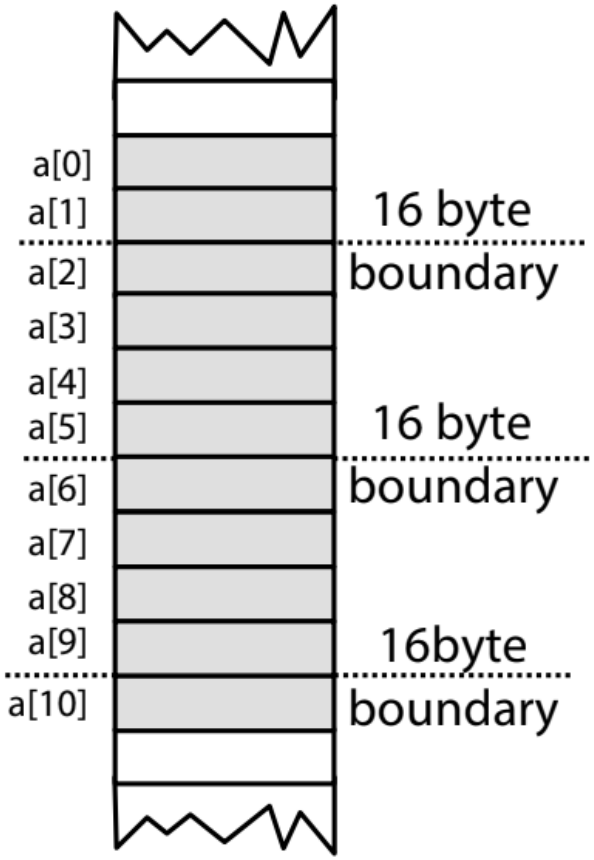


```
for (int i=0; i<11; i++)
    a[i] = 1.0;
```

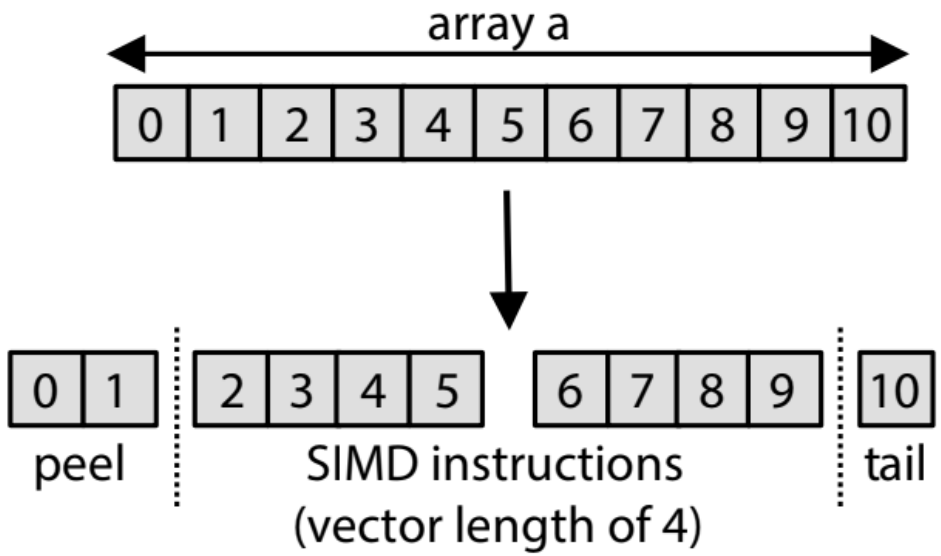

Gather and scatter scalar data elements in and out of a vector – Scalar data elements are gathered into vectors, operated on as a vector, and then scattered back to their destination locations.

```
for (int i=0; i<8; i+=2)
   a[i] = b[i] + c[i];
```

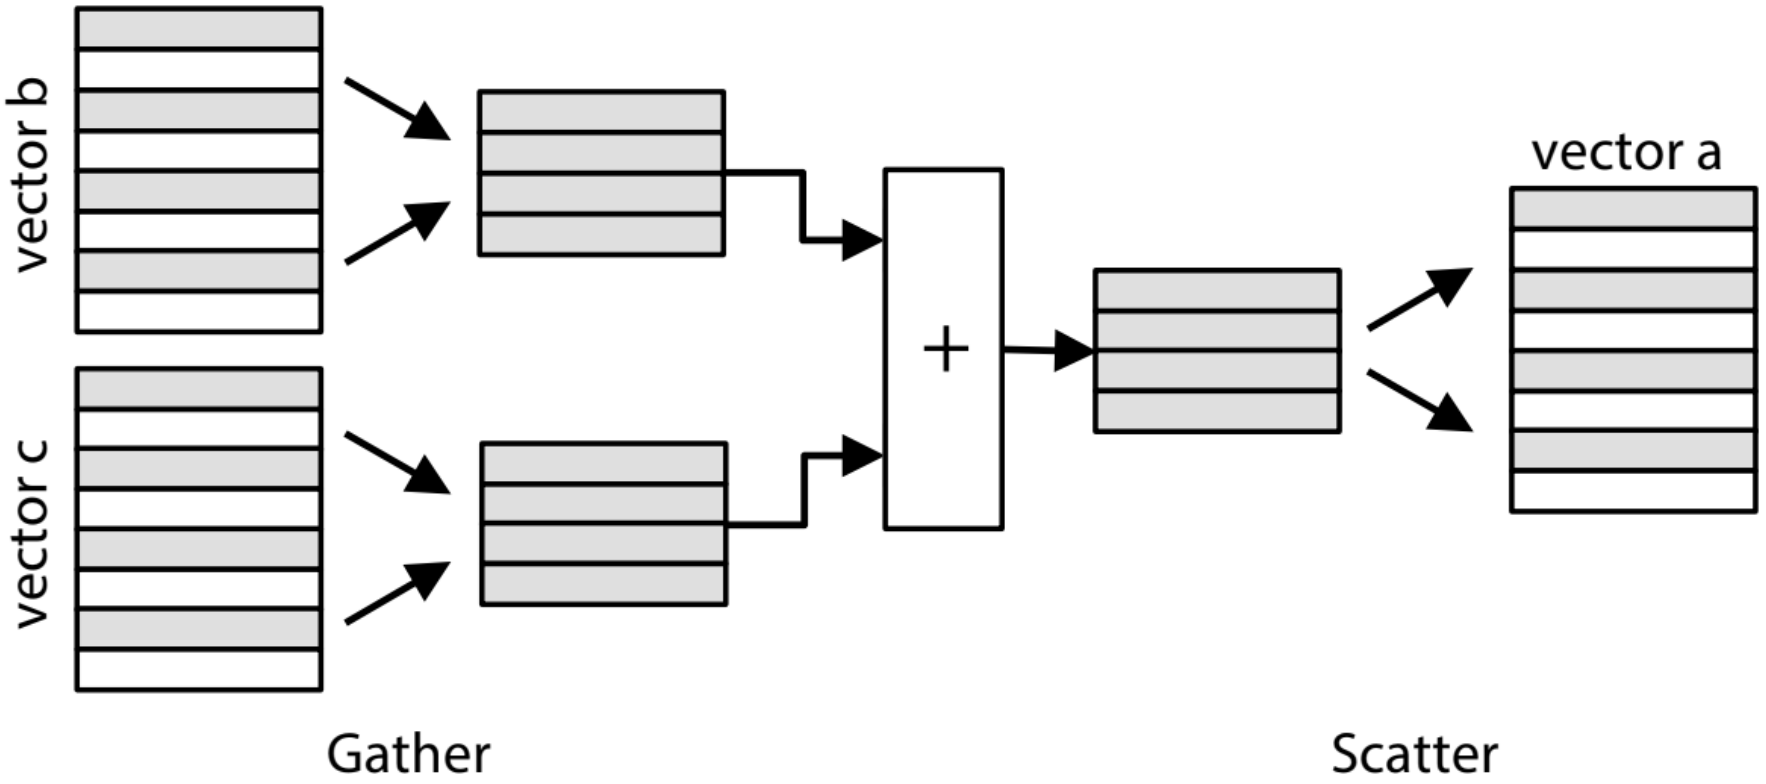

Combining thread and SIMD parallelism – Thread and SIMD parallelism are used to execute a loop.

```
#pragma omp for simd
for (int i=0; i<32; i++)
   a[i] += 1;
```

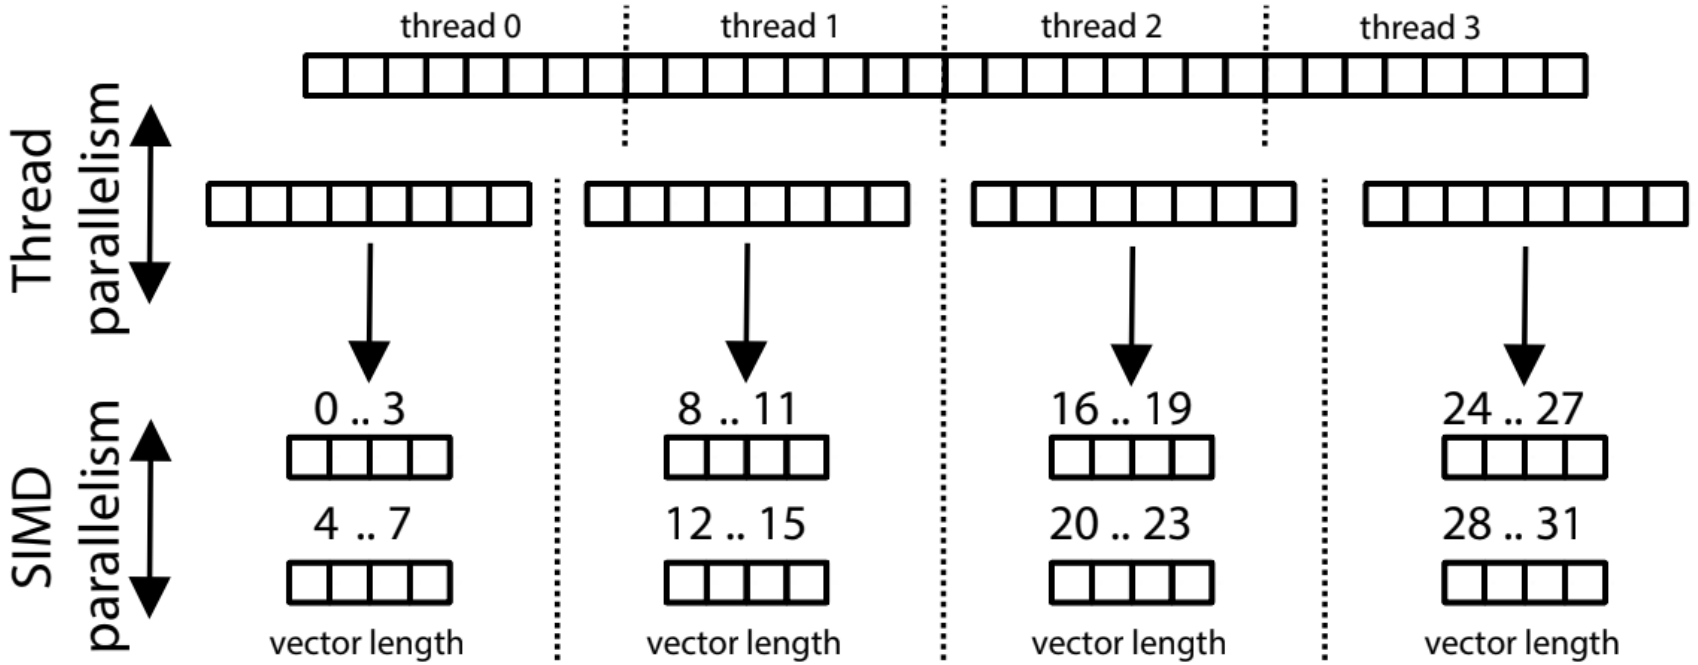

Illustration of function calls with SIMD – The scalar function f() is modified and renamed to F() in this example. This function supports vector input arguments and returns an entire vector.

```
for (int i=0; i<8; i++)
  a[i] = f(b[i]);
```

SIMD

```
for (int j=0; j<8; j+=4)
  a[0+j:4] = F(b[0+j:4]);
```

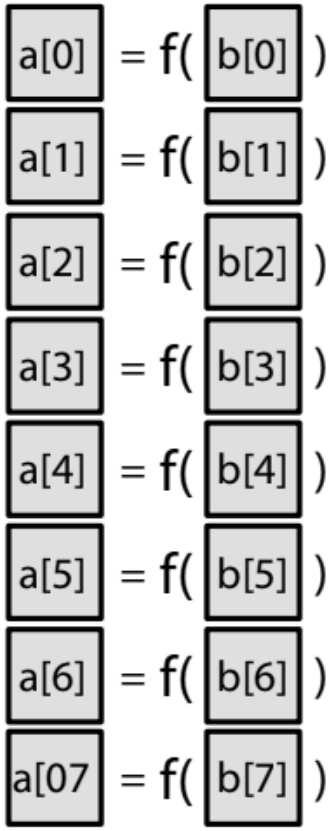

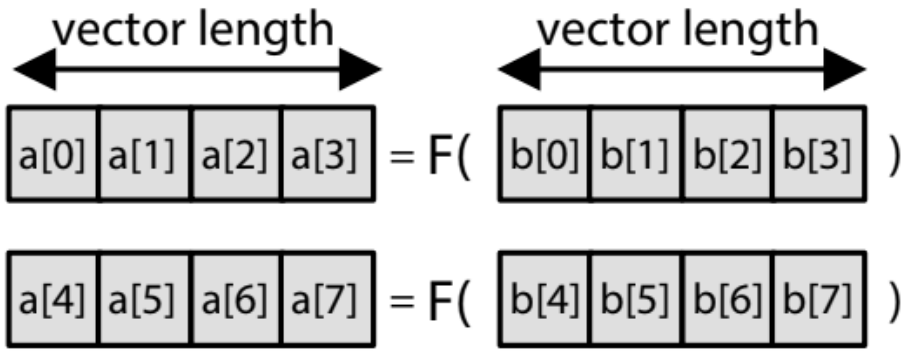

Conditional control flow converted to masked vector instructions – A vector mask predicate is used to enable or disable an operation on a given vector element. If the indexed mask is 1, the operation occurs. When it is 0, the operation is masked off.

```
for (int i=0; i<8; i++)
    if (b[i] < 0) a[i] += 1;
```